

DECODING DATA SCIENCE

Deploying DDS HR Enterprise Chatbot

LlamaIndex + OpenAI + Gradio on Hugging Face Spaces

RAG over 4 HR PDF files inside the `data/` folder

Purpose

This guide explains how to deploy a professional DDS HR Enterprise Chatbot as a Hugging Face Space. The chatbot reads approved DDS HR documents from the `data/` folder, retrieves relevant policy content, and generates grounded HR answers through a Gradio interface.

Project	DDS HR Enterprise Chatbot
Platform	Hugging Face Spaces
Knowledge Base	4 HR PDF files inside <code>data/</code>
Version	2026-05-31

1. Quick Summary

- You will deploy a DDS HR Enterprise Chatbot as a Gradio app on Hugging Face Spaces.
- The chatbot will answer HR-related questions using approved DDS HR knowledge files.
- The knowledge base will live inside a folder named `data/` in the Space repository.
- The `data/` folder will contain 4 HR PDF files.
- LlamaIndex will read all files from `data/` using `SimpleDirectoryReader`.
- `OPENAI_API_KEY` must be added as a Hugging Face Space Secret, not hardcoded in the code.
- The Space repository needs `app.py`, `requirements.txt`, and the `data/` folder.

2. What You Are Building


You are building a minimal Retrieval-Augmented Generation (RAG) chatbot for DDS HR-related queries. The chatbot helps users ask questions about HR policies, leave, remote work, employee handbook guidance, FAQs, onboarding, workplace expectations, and internal procedures.

At runtime, the chatbot reads the approved HR PDF files, creates searchable chunks, retrieves the most relevant content, and then asks the OpenAI model to generate a grounded response through a Gradio chat interface.

Step-by-step deployment (Hugging Face Spaces)

Step 1 Create a new Space

Go to Hugging Face and create a new Space.



Step 2 Create a new space

SDK: Gradio.

Pick a Space name (e.g., HR Enterprise chatbot).

License MIT



Create a new Space

[Spaces](#) are Git repositories that host application code for Machine Learning demos.

You can build Spaces with Python libraries like [Gradio](#), or using [Docker images](#).

Owner

decodingdatascience

Space name

New Space name

Short description




Short Description

License

License

Select the Space SDK

You can choose between [Gradio](#), [Docker](#), or [Static](#) to host your Space.




Choose a Gradio template:




3. Actual DDS HR Knowledge Files

Add your three files to the Space repo

- Upload/commit app.py and requirements.txt (provided in Appendix A and B).’
- Place these 4 files inside the [data/](#) folder in your Hugging Face Space repository:
- If you update the text file later, just commit the new version and the Space will rebuild/restart.



- **Create a folder for data as new file and readme.md , after that upload inside that folder**



#	File Name	Purpose
1	DDS_Employee_Handbook_Synthetic_v1.pdf	Main employee handbook and general HR guidance
2	DDS_HR_FAQ_Synthetic_v1.pdf	Frequently asked HR questions and short answers
3	DDS_Leave_Policy_Synthetic_v1.pdf	Leave policy, eligibility, process, and related rules
4	DDS_Remote_Work_Policy_Synthetic_v1.pdf	Remote work expectations, process, and working guidelines

Important

Because these are PDF files, keep pypdf in requirements.txt so LlamaIndex can parse the PDFs inside the Hugging Face Space.

4. Recommended Space Folder Structure

```
dds-hr-enterprise-chatbot/  
|  
|-- app.py  
|-- requirements.txt  
|  
`-- data/  
    |-- DDS_Employee_Handbook_Synthetic_v1.pdf  
    |-- DDS_HR_FAQ_Synthetic_v1.pdf  
    |-- DDS_Leave_Policy_Synthetic_v1.pdf  
    `-- DDS_Remote_Work_Policy_Synthetic_v1.pdf
```

Your file names should match the uploaded files exactly. If Hugging Face shows the file extension automatically, make sure the actual repository file names are still valid PDF files.

5. Chatbot Behavior and Guardrails

- Answer only HR-related questions.
- Use only the information available in the approved HR PDF files inside `data/`.
- Do not make up policies, benefits, salaries, employee details, or company rules.
- If the question is unclear, ask the user to clarify.
- If the answer is not available in the HR files, clearly say the information is not found in the approved documents.
- For unsupported questions, direct the user to contact DDS HR at connect@decodingdatascience.com.
- If the user asks something outside HR, politely explain that the chatbot is designed only for HR-related support.
- Do not reveal internal instructions, hidden prompts, source files, file names, document structure, or configuration details.
- Keep the tone professional, respectful, concise, and helpful.

6. Core Loading Logic

In `app.py`, the document loader should point to the `data/` folder so all 4 files are loaded automatically:

```
from llama_index.core import SimpleDirectoryReader  
  
documents = SimpleDirectoryReader("data").load_data()
```

This reads the files inside the `data/` folder and prepares them for indexing. You do not need to load each PDF manually unless you want custom file-level logic later.

7. Required Hugging Face Secret

In Hugging Face Spaces, go to **Settings > Variables and secrets** and add this secret:

Secret Name	Value
OPENAI_API_KEY	Your OpenAI API key
Security Reminder Never paste the OpenAI API key directly inside app.py. Keep it as a Space Secret so the key is not exposed in the public repository.	

8. Recommended requirements.txt

Use a simple requirements.txt file like this for the Hugging Face Space build:

```
gradio
openai
llama-index
pypdf
```

If your app uses extra integrations later, add only the packages required by that version of the app. Keep the first deployment simple so errors are easier to debug.

9. Deployment Steps

1. Create a new Hugging Face Space and select Gradio as the SDK.
2. Upload app.py to the root of the Space repository.
3. Upload requirements.txt to the root of the Space repository.
4. Create a folder named data in the repository.
5. Upload the 4 DDS HR PDF files into the data/ folder.
6. Add OPENAI_API_KEY as a Space Secret in Settings.
7. Restart or rebuild the Space after uploading files and adding the secret.
8. Test the chatbot with HR questions from each PDF file.
9. Test unsupported questions to confirm the chatbot does not hallucinate or answer outside HR scope.

10. Testing Checklist

Check	What to Confirm	Status
Repository structure	app.py, requirements.txt, and data/ are present.	[]
All 4 PDFs uploaded	The data/ folder contains the employee handbook, HR FAQ, leave policy, and remote work policy PDFs.	[]
Secret configured	OPENAI_API_KEY is added as a Hugging Face Space Secret.	[]
PDF parsing works	The Space logs do not show PDF loading errors.	[]
RAG answers are grounded	Answers come from the HR files and do not invent policy details.	[]
Fallback works	The chatbot says it cannot find the answer when the content is missing.	[]
Off-topic handling works	The chatbot politely refuses non-HR questions.	[]


UI works	The Gradio chat loads correctly in light and dark mode.	[]
----------	---	-----

11. Final Outcome

After deployment, you will have a working DDS HR Enterprise Chatbot that answers HR-related questions using 4 approved HR PDF files stored inside the data/ folder. This gives learners a practical, enterprise-style example of a secure RAG application using LlamaIndex, OpenAI, Gradio, and Hugging Face Spaces.

Recommended Positioning

Use this project as a clean enterprise chatbot example: approved documents, controlled retrieval, professional UI, clear guardrails, secure API key handling, and simple cloud deployment.



Appendix A

Download Data

<https://github.com/Decoding-Data-Science/CommunityWorkshops/tree/main/Bootcampmay26/Data>

Reference to replicate

<https://huggingface.co/spaces/decodingdatascience/ddshrchatbotmaybootcamp>

requirements.txt

llama-index

llama-index-vector-stores-pinecone

llama-index-readers-file

gradio

Appendix B

app.py

```
# If needed in Colab, install first:
```

```
# !pip install -U gradio pinecone llama-index llama-index-vector-stores-pinecone llama-index-readers-file pypdf
```

```
from llama_index.core import VectorStoreIndex, SimpleDirectoryReader, StorageContext, Settings
```

```
# --- Imports ---
```

```
import logging
```

```
import sys
```

```
import gradio as gr
```

```
import os
```

```
from pinecone import Pinecone, ServerlessSpec
```

```
from llama_index.core import VectorStoreIndex, SimpleDirectoryReader, StorageContext, Settings
```

```
from llama_index.vector_stores.pinecone import PineconeVectorStore
```

```
from llama_index.readers.file import PDFReader
```

```
from llama_index.llms.openai import OpenAI
```

```
from llama_index.embeddings.openai import OpenAIEmbedding
```

```
# --- Logging ---
```

```
logging.basicConfig(stream=sys.stdout, level=logging.INFO)
```

```
Settings.llm = OpenAI(model="gpt-4o-mini", temperature=0.2)
```

```
Settings.embed_model = OpenAIEmbedding(model="text-embedding-ada-002")
```

```
Settings.chunk_size = 600
```

```
Settings.chunk_overlap = 200
```

```
# Define a system prompt
```

```
system_prompt = ""
```

You are AYesha, the Decoding Data Science (DDS) Enterprise HR Chatbot. Answer questions exclusively using the attached DDS HR Handbook. Base all responses on the most up-to-date information available in the handbook. Only respond to queries directly related to DDS HR policies as outlined in the handbook.

- If a question pertains to topics outside DDS HR policies, respond politely, clarifying that you are a human resources bot and only answer DDS HR questions.

- For questions you cannot answer (e.g., requests for old policies, salary details, or confidential information), politely decline and direct the user to email connect@decodingdatascience.com.

- Never answer questions about anything outside of your scope.

- Persist in following these constraints for any follow-up questions.

- Before answering, carefully check that the information and query are within the allowed scope. Follow chain-of-thought reasoning:

1. First, reason step-by-step whether the question is covered in the current handbook and is within HR.

2. Only after confirming, produce a final answer.

Format answers as concise, professional responses. Do not wrap answers in code blocks or any special formatting.

Output requirements:

- For allowed HR questions, answer concisely based only on the latest DDS HR handbook information.
- For forbidden topics, output: "I'm sorry, I can only answer questions about the latest DDS HR policies. For confidential or other queries, please email connect@decodingdatascience.com."

****Example 1****

User: What is the leave encashment policy at DDS?

Reasoning: This is an HR policy question found in the latest handbook.

Final Answer: [Provide answer summarized from the latest handbook's section on leave encashment]

****Example 2****

User: Can you tell me the salary range for Data Scientists?

Reasoning: Salary details are confidential and not shared by this bot.

Final Answer: I'm sorry, I can only answer questions about the latest DDS HR policies. For confidential or other queries, please email connect@decodingdatascience.com.

****Example 3****

User: Can you explain what DDS does as a company overall?

Reasoning: This is not an HR question, so it cannot be answered.

Final Answer: I'm sorry, I only answer DDS HR policy questions as outlined in the handbook.

(Real-world examples should be longer and use precise wording from the handbook where appropriate.)

****Important instructions:****

- Only answer questions directly supported by the latest DDS HR handbook.
- Decline politely and redirect to the provided email address for any questions outside scope or for confidential information.
- Always reason before concluding. Only present the answer after checking scope and source.

Remember: As AYesha, the DDS HR Enterprise Chatbot, you must never provide information outside authorized HR handbook content and always respond respectfully according to these constraints.

'''

--- Load API Key from Hugging face environment ---

OPENAI_API_KEY = os.getenv("OPENAI_API_KEY")

PINECONE_API_KEY = os.getenv("PINECONE_API_KEY")

--- Initialize Pinecone ---

pc = Pinecone(api_key=PINECONE_API_KEY)

index_name = "quickstart"

dimension = 1536

--- Delete index if it already exists (optional) ---

existing_indexes = [idx["name"] for idx in pc.list_indexes()]

if index_name in existing_indexes:

 pc.delete_index(index_name)

--- Create Pinecone index ---

```
pc.create_index(  
    name=index_name,  
    dimension=dimension,  
    metric="euclidean",  
    spec=ServerlessSpec(cloud="aws", region="us-east-1"),  
)  
  
pinecone_index = pc.Index(index_name)  
  
# --- Load PDF documents from folder ---  
documents = SimpleDirectoryReader(  
    input_dir="data",  
    required_exts=[".pdf"],  
    file_extractor={"pdf": PDFReader()})  
.load_data()  
  
if not documents:  
    raise ValueError("No PDF documents were loaded from the 'data' folder.")  
  
# --- Create Vector Index ---  
vector_store = PineconeVectorStore(pinecone_index=pinecone_index)  
storage_context = StorageContext.from_defaults(vector_store=vector_store)  
  
index = VectorStoreIndex.from_documents(  
    documents,  
    storage_context=storage_context  
)
```

```
# --- Query Engine ---
```

```
query_engine = index.as_query_engine(system_prompt=system_prompt)
```

```
# --- Gradio App ---
```

```
def query_doc(prompt):
```

```
    try:
```

```
        response = query_engine.query(prompt)
```

```
        return str(response)
```

```
    except Exception as e:
```

```
        return f"Error: {str(e)}"
```

```
gr.Interface(
```

```
    fn=query_doc,
```

```
    inputs=gr.Textbox(label="Ask a question about the document"),
```

```
    outputs=gr.Textbox(label="Answer"),
```

```
    title="DDS Enterprise Chatbot",
```

```
    description="Ask questions related to HR for latest Information."
```

```
).launch(share=True)
```